

Théo Matricon

2025–current

Postdoc with Mathieu Acher

Code Generation, Reproducibility

DiverSE, Software Engineering team

IRISA, Rennes

2021–2024

PhD supervised by Nathanaël Fijalkow


Scaling domain agnostic techniques for program synthesis

M2F, Formal Methods team

LaBRI, Bordeaux

Medical Condition (RQTH)


Research Contributions



Selected Publications

- 2026, TOSEM (Q1) + another under review
- 2026, MSR (A) + Distinguished Paper Award
- 2025, TACAS (A)
- 2025, AAI (top 20%) (A*)
- 2025, IST (Q1)
Postdoc
- 2022, AAI (top 20%) (A*)

PhD Motivation



PhD Motivation: In Practice

Specifications

Logic

$\forall a, b$
 $f(a, b) \geq a$
 $f(a, b) \geq b$
 $f(a, b) \in \{a, b\}$

Examples

$f(1, 5) = 5$
 $f(2, 1) = 2$
 $f(-3, -9) = -3$

Natural Language

‘Write a function that takes the maximum of its two arguments.’

Produced Algorithm

```
def max(a: int, b: int) ->int:
    if a <=b:
        return b
    else:
        return a
```

Program Synthesis: Problem


Input

- the search space G :
a deterministic tree grammar
- a specification $\mathcal{C}$:
checks if a program $p \in \mathcal{L}(G)$ matches the specification


Output

- a program in the search space that matches the specification:
 $p \in \mathcal{L}(G)$ such that $\mathcal{C}(p) = \checkmark$

Program Synthesis: Frameworks



Program Synthesis: Frameworks



Program Synthesis: Enumeration Problem

Input

- the search space G :
a deterministic tree grammar with a cost for each tree
- a specification $\mathcal{C}$:
checks if a program $p \in \mathcal{L}(G)$ matches the specification

Goal


- enumeration of programs in order of non-decreasing costs

Delay

- time complexity between enumeration of two successive programs in terms of number of enumerated programs

Enumeration Algorithm Before Our Contribution

priority queue: S pairs of (program, cost)




At step n

$$S = O(n)$$

$$k = O(1)$$

Our Contribution: Constant Delay

bucket queue: M buckets of (programs, cost)



At step n


$$M = O(1)$$

$$k = O(1)$$

Our Contribution: Empirical Results

Major Papers

- 2017, ML + **Enum.**, Balog *et al.*, ICLR
- 2018, $O(\log n)$, Lee *et al.*, PLDI



Our Contribution: Summary

We prove bounded differences in cost:


$$\exists M, \forall n, \text{cost_next}(p_n) - \text{cost}(p_n) \leq M$$

This implies: **priority queues** $O(\log n)$ $\rightarrow$ **bucket queues** $O(1)$


Impact

Published in **AAAI 2025** (+oral: 20% of accepted papers).
Fastest ranked enumeration for program synthesis **in practice**.
First algorithm with $O(1)$ delay $\rightarrow$ closes open question.

Evaluations



Evaluations: In Practice



My Research Project

Observations

- Exponential growth in scientific publications [1]
- Reproducibility crisis in computational science
- Evaluations are a key aspect of reproducibility

My Research Project

Observations

- **Exponential growth** in scientific publications [1]
- **Reproducibility** crisis in computational science
- **Evaluations** are a key aspect of **reproducibility**

How to design **reliable and **efficient** evaluations?**

Reliable → formal guarantees

Efficient → reduce cost of **evaluations**

My Research Project

Observations

- **Exponential growth** in scientific publications [1]
- **Reproducibility** crisis in computational science
- **Evaluations** are a key aspect of **reproducibility**

How to design **reliable** and **efficient evaluations**?

Reliable → formal guarantees

Efficient → reduce cost of **evaluations**

Main Challenges

SE Modelling the huge hardware and software **variability**

AI Predicting heterogeneous **evaluation** data

FM Giving formal guarantees while still being **efficient**

My Research Project

Current — no formal guarantees · costly · hard to reproduce

Candidates
+ Benchmark

Run **everything**

$\hat{r}$

My Research Project

Current — no formal guarantees · costly · hard to reproduce

Candidates
+ Benchmark

Run **everything**

$\hat{r}$

My Vision — reliable · efficient · reproducible

Candidates
+ Benchmark

My Research Project

Current — no formal guarantees · costly · hard to reproduce

Candidates
+ Benchmark

Run **everything**

$\hat{r}$

My Vision — reliable · efficient · reproducible

Candidates
+ Benchmark

Axis 1: Reliable
Bayesian U.Q.

select

observe

guarantee

Axis 2: Efficient
Adaptive eval.

stop?

Evaluate
(subset)

My Research Project

Current — no formal guarantees · costly · hard to reproduce

Candidates
+ Benchmark

Run **everything**

$\hat{r}$

My Vision — reliable · efficient · reproducible

Candidates
+ Benchmark

Axis 1: Reliable
Bayesian U.Q.

select

observe

guarantee

Axis 2: Efficient
Adaptive eval.

stop?

Evaluate
(subset)

$$\mathbb{P}(r \in [\hat{r} - \Delta; \hat{r} + \Delta]) \geq 1 - \alpha$$

Classic Approach vs. Our Approach

Classic Approach

What is evaluated

Run all n instances — exhaustively

Our Approach

What is evaluated

Adaptive subset + early stopping
(~50% fewer runs [1])

Classic Approach vs. Our Approach

Classic Approach

What is evaluated

Run all n instances — exhaustively

Statistical framework

assumes i.i.d. & normality

Our Approach

What is evaluated

Adaptive subset + early stopping
(~50% fewer runs [1])

Statistical framework

Adjustable complexity

Classic Approach vs. Our Approach

Classic Approach

What is evaluated

Run all n instances — exhaustively

Statistical framework

assumes i.i.d. & normality

Variability

Single environment assumed
results may not transfer

Our Approach

What is evaluated

Adaptive subset + early stopping
(~50% fewer runs [1])

Statistical framework

Adjustable complexity

Variability

Multi-environment hierarchical model
quantifies generalizability

Classic Approach vs. Our Approach

Classic Approach

What is evaluated

Run all n instances — exhaustively

Statistical framework

assumes i.i.d. & normality

Variability

Single environment assumed
results may not transfer

Reproducibility

binary: reproduced or not

Our Approach

What is evaluated

Adaptive subset + early stopping
(~50% fewer runs [1])

Statistical framework

Adjustable complexity

Variability

Multi-environment hierarchical model
quantifies generalizability

Reproducibility

graded: exact → ordinal → trend

Axis 1: Reliable Evaluations

Variability

CPU-A • GCC-12 • Linux

$$\hat{r}_A = 0.82 \quad \hat{r}_B = 0.75 \quad A \succ B$$

CPU-B • GCC-14 • Linux


$$\hat{r}_A = 0.79 \quad \hat{r}_B = 0.77 \quad A \succ B$$

CPU-C • Clang • macOS


$$\hat{r}_A = 0.71 \quad \hat{r}_B = 0.74 \quad A \prec B$$

inconsistent?

Axis 1: Reliable Evaluations




Axis 1: Reliable Evaluations



Axis 2: Efficient Evaluations

Execution on an instance is not atomic



Integration

Transversal project with many applications in empirical fields

→ **CRIStAL**, UMR 9189, Lille, **Spirals** (SEEDS)

Frugality and Evaluations: Clément Quinton, Romain Rouvoy

Trust: Pierre Bourhis (Citadelle)


Statistics: Scool team

→ **IRISA**, UMR 6074, Rennes, **DiverSE** (rRaise)

Reproducibility: Mathieu Acher, Olivier Barais, Djamel E. Khelladi


Frugality: *Romain Lefeuvre, Quentin Perez*

Théo Matricon



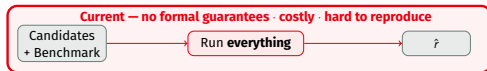


Selected Publications

2 AAAI →
IST, MSR, 2 TOSEM
TACAS



Research Project



- [1] G. Bathie, N. Fijalkow, T. Matricon, B. Mouillon, and P. Vandenhove, "LTL_f learning meets boolean synthesis," in *Tools and Algorithms for the Construction and Analysis of Systems: 32nd International Conference, TACAS 2026, Held as Part of the International Joint Conferences on Theory and Practice of Software, ETAPS 2026, Proceedings*, 2026. [Online]. Available: <https://theomat.github.io/files/ltl.pdf>
- [2] R. Robbes, T. Matricon, T. Degueule, A. Hora, and S. Zacchiroli, "Promises, perils, and (timely) heuristics for mining coding agent activity," in *2026 IEEE/ACM 23rd International Conference on Mining Software Repositories (MSR)*, IEEE, 2026.
- [3] A. Martin, D. E. Khelladi, T. Matricon, and M. Acher, "Re-evaluating metamorphic testing of chess engines: A replication study," *Information and Software Technology*, vol. 181, p. 107 679, 2025, ISSN: 0950-5849. DOI: <https://doi.org/10.1016/j.infsof.2025.107679> [Online].

Available: <https://www.sciencedirect.com/science/article/pii/S0950584925000187>

- [4] T. Matricon, M. Acher, H. Spieker, and A. Gotlieb, “Efficiently ranking software variants with minimal benchmarks,” *arXiv preprint arXiv:2509.06716, under revision at TOSEM*, 2025.
- [5] T. Matricon, N. Fijalkow, and G. Lagarde, “Ecosearch: A constant-delay best-first search algorithm for program synthesis,” in *Proceedings (AAAI Artificial Intelligence and Interactive Digital Entertainment Conference)*, 2025. [Online]. Available: <https://arxiv.org/abs/2412.17330>
- [6] R. Robbes, T. Matricon, T. Degueule, A. Hora, and S. Zacchiroli, “Empirical study of coding agent adoption on github,” *Under submission*, 2025.
- [7] G. Shabadi, N. Fijalkow, and T. Matricon, “Theoretical foundations for programmatic reinforcement learning,” in *GenPlanAI workshop AAAI (preprint)*, 2025. [Online]. Available: <https://doi.org/10.48550/arXiv.2402.11650>

- [8] H. Spieker, T. Matricon, N. Belmecheri, J. E. Betten, G. L. B. Lyan, H. Borges, Q. Mazouni, D. Gross, A. Gotlieb, and M. Acher, "Prompting for performance: Exploring llms for configuring software," in *The 37th IEEE International Conference on Tools with Artificial Intelligence (ICTAI)*, 2025. [Online]. Available: <https://arxiv.org/pdf/2507.09790>
- [9] T. Matricon, N. Fijalkow, and G. Margueritte, "Wikicoder: Learning to write knowledge-powered code," in *Model Checking Software*, G. Caltais and C. Schilling, Eds., Springer Nature Switzerland, 2023, pp. 123–140, ISBN: 978-3-031-32157-3. [Online]. Available: <https://rdcu.be/dITGA>
- [10] M. Anastacio, T. Matricon, and H. Hoos, "Challenges of acquiring compositional inductive biases via meta-learning," in *ECMLPKDD Workshop on Meta-Knowledge Transfer*, P. Brazdil, J. N. van Rijn, H. Gouk, and F. Mohr, Eds., ser. Proceedings of Machine Learning Research, vol. 191, PMLR, Sep. 2022, pp. 11–23. [Online]. Available: <https://proceedings.mlr.press/v191/anastacio22a.html>

- [11] N. Fijalkow, G. Lagarde, T. Matricon, K. Ellis, P. Ohlmann, and A. Potta, “Scaling neural program synthesis with distribution-based search,” in *International Conference on Artificial Intelligence, AAAI*, 2022. [Online]. Available: <https://arxiv.org/abs/2110.12485>
- [12] T. Matricon, N. Fijalkow, G. Lagarde, and K. Ellis, “Deepsynth: Scaling neural program synthesis with distribution-based search,” *Journal of Open Source Software*, vol. 7, no. 78, p. 4151, 2022. DOI: 10.21105/joss.04151
- [13] T. Matricon, M. Anastacio, N. Fijalkow, L. Simon, and H. H. Hoos, “Statistical Comparison of Algorithm Performance Through Instance Selection,” in *27th International Conference on Principles and Practice of Constraint Programming (CP 2021)*, L. D. Michel, Ed., ser. Leibniz International Proceedings in Informatics (LIPIcs), vol. 210, Dagstuhl, Germany: Schloss Dagstuhl – Leibniz-Zentrum für Informatik, 2021, 43:1–43:21, ISBN: 978-3-95977-211-2. DOI:

10.4230/LIPIcs.CP.2021.43 [Online]. Available: <https://drops.dagstuhl.de/entities/document/10.4230/LIPIcs.CP.2021.43>

- [14] D. Duplyakin, A. Uta, A. Maricq, and R. Ricci, "In datacenter performance, the only constant is change," in *2020 20th IEEE/ACM International Symposium on Cluster, Cloud and Internet Computing (CCGRID)*, IEEE, 2020, pp. 370–379.
- [15] N. A. Ernst, "Bayesian hierarchical modelling for tailoring metric thresholds," in *Proceedings of the 15th international conference on mining software repositories*, 2018, pp. 587–591.
- [16] L. Bornmann and R. Mutz, "Growth rates of modern science: A bibliometric analysis based on the number of publications and cited references," *Journal of the association for information science and technology*, vol. 66, no. 11, pp. 2215–2222, 2015.